

# Software Project Report: LATTICE WATERING

VALENTIN PICKEL, JONAS HEINEMANN, CHRISTIAN MÜLLER, and KAAAN DÖNMEZ, Freie Universität Berlin, Germany

## 1 Introduction

We may start from scratch by reminiscing the the first meeting of all course participants on April 4th. This software project course should dive into the ecosystem internet. To do so, a rough split was set between the worlds of the *Internet of Things* (IoT), which deals with low power devices in large networks plagued with packet loss, and the conventional internet and its vastly more potent desktop computers and servers. Especially, us students should become aware of design decisions with respect to the communication solutions that our software should utilize. The project *Lattice Watering* eyes out the first of both worlds.

### 1.1 Idea Outline

After suggesting the building an IoT application for watering plants automatically on May 2nd, the software project *Lattice Watering* and its group formed on the 9th of May, comprising of the aforementioned members. The term *Lattice* is a reference to the studies of crystal/grid-like molecular structures in physics settings, which are often represented by two-dimensional structures called lattices. The first idea was not just to automate the watering of the plants, but also to let the boards drive around on a large grid. As it turns out, the hardware requirements for such a project would be too difficult to implement, so after some discussion we stuck to the original idea of a plant watering application, not changing the initial name.

In a nutshell, the idea is to be able to use a normal desktop computer to control several smaller microcontrollers, that are scattered widely. This may happen by setting automatic rules, or by the microcontrollers automatically turning on the water pumps as soon as humidity levels, which are measured by sensors, drop beyond a certain threshold. The latter aspect can be configured in a web frontend that is designed by us. Besides watering the plants, we also collect statistics like the humidity levels or network usage.

### 1.2 Team Organization and Applied Practices

For a detailed overview of the code hierarchy, one can look into the file `README.md` inside of the project directory. *git* was used as our version control system, with a repository setup on *GitHub*. Our team members usually did not have fixed tasks, but utilize the agile method of *Kanban* with the popular free software license by the *MIT*. On *GitHub*, a Kanban board, see fig. 1, was setup, in which members can add new issues and tasks to three categories: *Backlog*, *In Progress* and *Done*. This allowed our work to proceed in a assembly line-type way, in the sense that there are certain tasks that are saved up for

later. However, in the nature of things, some members prefer to work closer to the IoT devices themselves, whilst others work e.g. with the web frontend. It is important to mention that we used *agile* methods, as we all had different experiences in the field of IoT communication we were studying, so agile methods and rapid prototyping seemed like good methods for this project.

We consider it noteworthy that for uniformity in code, we set up some *coding conventions*, which are described in the `README.md` file. Such include not just whether to use the standardized integer types from `inttypes.h` in C, but also the use of formatting tools and static analyzers. Although they are quite few in this project, we hope they reflect that we tried take care of code quality at all times, and using matured automatic tools for that eases a normally administrative task. But we also missed out on a few things. Like developing many more tests and using continuous integration. But at this scale these tools did not seem necessary anyways.

Communication was done via the proprietary communication platform *Discord*, which was preferred over a decentralized solution like *Element/Matrix* due to all course members already using it. We may improve on that in a future project.

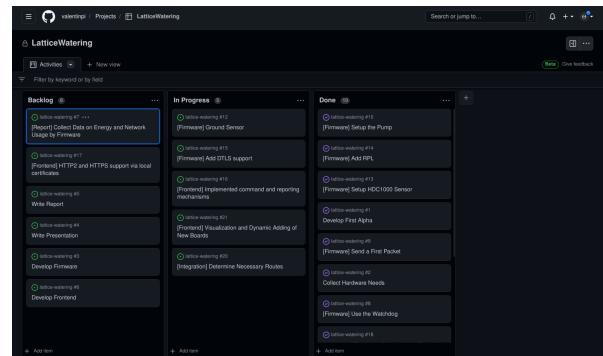


Fig. 1. The kanban project board. (State: 13th June, 18:32)

## 2 Hardware

One major motivation for us was to work close to hardware. The boards and sensory we used came mostly from suggestions by Hauke Petersen, a scientific coworker in the department and lead developer of RIOT OS. In our network, see section 2, we essentially use three types of devices: A host computer which is quite powerful, a border router that translates traffic of a SixLoWPAN network, that is the name for a set of protocols used in low power radio based communication networks, and so-called nodes. The latter two types of boards use class 2 embedded devices with about 32kB of SRAM memory. A sharp eye may criticize that the border router is also quite a

Authors' address: Valentin Pickel, valentinpickel@gmx.de; Jonas Heinemann; Christian Müller; Kaan Dönmez, Institute for Computer Science, Freie Universität Berlin, Berlin, Germany.

weak board, but during development we tried to mitigate this possible problem by reducing its role in the communication network to a pure routing one. In total, we have the following hardware list:

- (1) One personal computer with the software set up.
- (2) Two SAMR21-XPRO boards, one border router and one node.
- (3) Two micro USB cables, optionally a battery connector or powerbank.
- (4) One DRV8833 motor driver board.
- (5) One electronic pump, see the file mentioned below. It should come with a long tube.
- (6) One capacitive moisture sensor, see the file mentioned below.
- (7) Nine female jumper cables.
- (8) Five male jumper cables.

This entire list can be found in `HWSETUP.md` inside of the repository, with descriptions on how to connect the jumper cables. In the repository, one can also find instructions on how to setup the software. However, we assume that the reader is already familiar with RIOTOS and has the appropriate toolchains and udev rules set up.

As for the hardware itself, we want to share only one more story during development. There was a very specific problem at the start of the project, which halted development shortly, but we managed to catch up as we were quite early when we tried to obtain hardware. The problem is that connecting the pump to the node boards themselves, the ones of model SAMR21-XPRO, the pump induces a current of 200mA. This is too much for the GPIO pins to handle, as one can read in the manual for the board, which we have provided in the `man` folder of the repository.

Together with Hauke, we discussed possible solutions to this problem, of which at first the most promising one can be seen in fig. 2. We can see that in this schematic the pump is connected to the board via the GPIO connection and both the 5 volt and ground pins. We shall not go into the details of the electronics of that circuit and why it would work under ideal conditions. According to Hauke this is rather a matter of experience. What the may note is that the essential idea is that current flows between the 5V0 and GND pins to drive the pump if and only if as soon as an electric signal on the PB03 pin is raised and thus the transistor is switched. We built the circuit, but as it turns out the transistor we had available failed. So Hauke proposed that we obtain some predesigned motor boards, which are documented in `HWSETUP.md`. Those boards already have this circuit idea embedded in one board, meaning that there are some resistors and a transistor built onto one board and we just have to connecting everything properly.

To get a feeling for what an almost fully connected node looks like, one may refer to fig. 3.

### 3 Network Structure

#### 4 Application Structure

In total, we have developed *five* different applications which all work together.

- `fw`: The firmware that is deployed on the nodes themselves reads out sensor values and sends them the host.
- `br`: The border router acts as a mediator between the lossy SixLoWPAN and the local Ethos network with the host. It allows the nodes that are reachable to send messages to the host computer.
- `proxy`: The proxy receives DTLS traffic from the nodes, decrypts them, appends a target IP and forwards them to the backend software written in Node.JS.
- `front`: The backend software serves the webpage for the frontend and manages the database.
- `www`: The website written in JavaScript, which may or may not be considered a separate application, offers the user

To enable DTLS for our application, we had to first include the `gcoap_dtls` package in RIOT with a PRNG (Pseudo Random Number Generator) and an appropriate library, in this case `tinydtls`. However, Node.JS support for DTLS is in quite a rough state as this time, as none of the libraries we were trying out worked with. To be precise, we tried the following libraries:

- `dtls` [1]: This NPM package just does not contain anything.
- `openssl-dtls` [2]: Node FFI bindings that work, but during development, OpenSSL did not work with the ciphersuite `TLS_PSK_WITH_AES_128_CCM_8`, which is used by `tinydtls`. The cipher negotiation fails for to us unknown reasons.
- `node-mbed-dtls` [3]: Based on `mbed-dtls`, but with the same issue as with `openssl-dtls`.
- `werift-dtls` [4]: Crashes upon receival of a Client Hello packet, also very poorly documented code. As it seems, it is part of a bigger set of tools for WebRTC.
- `nodejs-dtls` [5]: Only contains a DTLS client.
- `@nodertc/dtls` [6]: Only contains a DTLS client. Development for a server seems to be ongoing on the main branch on GitHub, but with no recent activity
- `node-dtls-proxy` [7]: Crashes upon receival of a Client Hello packet.
- `goldy` [8]: Cipher negotiation with a node fails, also it seems that the proxy only forwards traffic between two peers, but we want to be dynamic in that multiple nodes can connect to the server and we do not need to start a proxy for each one separately.

There was also a short unsuccessful attempt at writing a very small proxy using `python-dtls`, but this was conceptually flawed, as the proxy needs to know where to route the packets from the frontend. In the end, as the board uses `tinydtls`, we decided to use Rust to write a small proxy using `tinydtls-sys`. ?? shows the structure of the proxy, which acts in the middle of the communication between the frontend and the nodes. Note that the proxy is running on the host computer.




Fig. 2. Original circuit idea for the pump. This is a rough sketch, and does not exactly represent the circuitry involved. (P,S,G,D) = (Pump, Source, Gate, Drain). We credit Hauke Petersen with the design of this circuit.




Fig. 3. An almost fully connected node. We can see that it is quite a cable salad, in the final presentation we gave we suggested to put everything in a casing so that it looks more compact. On the right we can see the driver board that connects the pump and the board, and on the left the capacitive moisture sensor that one would put into soil to measure humidity. It is not *fully* connected, because for one we would place the moisture sensor into soil, the pump into water, whilst connecting it with the plastic tube shown on the right, and then we would connect the board to a power source. This could be a battery connector, here the large case on the right that can hold up to four AA batteries, or a simple micro USB cable.

When a board wants to send data through the `/data` CoAP POST method of the host, it first tries to establish a secure DTLS session. For that, it sends a Client Hello packet and then does cipher negotiation such that the future communication partners know how to encrypt and decrypt packages. One special property of our communication is that we use *Pre-Shared Keys* (PSKs). Via an *identity message*, the server tells the node which pre-shared key to use. We only have one, so we called our identity default. The PSKs themselves are just random data, here 16 bytes, as RIOT does not support

longer keys. We automatically generate them using a script, and recommend that inside of `README.md`.

## 5 Closing Words

### References

- [1] dtls. NPM Package: <https://www.npmjs.com/package/dtls> [last access 25th July 18:37].
- [2] openssl-dtls. NPM Package: <https://www.npmjs.com/package/openssl-dtls> [last access 25th July 18:45].
- [3] node-mbed-dtls. NPM Package: <https://www.npmjs.com/package/node-mbed-dtls> [last access 25th July 19:32].
- [4] werift-dtls. NPM Package: <https://www.npmjs.com/package/werift-dtls> [last access 25th July 19:33].
- [5] nodejs-dtls. NPM Package: <https://www.npmjs.com/package/nodejs-dtls> [last access 25th July 19:35].
- [6] @nodertc/dtls. NPM Package: <https://www.npmjs.com/package/@nodertc/dtls> [last access 25th July 19:41].
- [7] node-dtls-proxy. NPM Package: <https://www.npmjs.com/package/node-dtls-proxy> [last access 25th July 19:43].
- [8] goldy. GitHub Repository: <https://github.com/ibm-security-innovation/goldy> [last access 25th July 19:44].




Fig. 4. The Network Topology. The device interfaces are named after our code. The `tap0` device is named after the typical Linux entry. The network stack is based on COAP, DTLS, UDP, 6LoWPAN IPHC and FRAG, RPL, IPv6. The global IP addresses of the nodes inside the `fc00:0:0:0:0:0:0:1::/64` network are chosen by appending the Layer 2 addresses to the network prefix. The DAG structure of the RPL network is shown schematically.

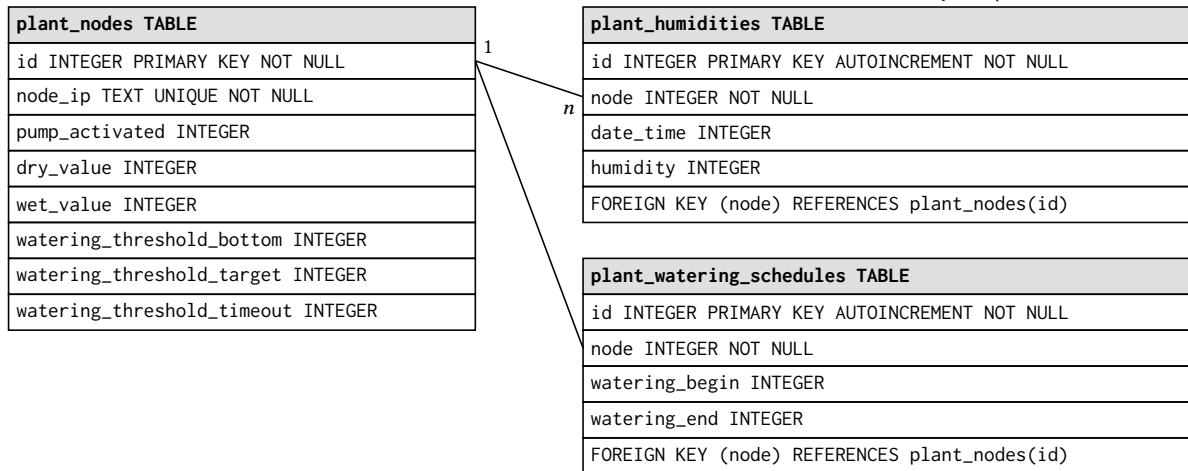


Fig. 6. The SQLite database schema. The left table persists the current configurations for each node. The network structure guarantees the uniqueness of the IPs, so it suffices to use them as an identifier in our use case. One may consider replacing this with an application-specific ID which would have to persist on the nodes themselves, e.g. via EEPROM. One can find this exact structure in `front/db.js`.

